

stages:

- build
- test
- deploy

build_job:

stage: build

script:

- echo "Compiling the code..."

unit_tests:

stage: test

script:

- echo "Running unit tests..."

integration_tests:

stage: test

script:

- echo "Running integration tests..."

deploy_job:

stage: deploy

script:

- echo "Deploying to production..."

Jobs inside the same stage run in parallel (concurrently), **as long as enough GitLab Runners are available**.

In the above example:

- unit_tests and integration_tests belong to the **test stage**.
- They run **at the same time** on different runners.
- This reduces overall pipeline time.

Stage Execution Order

- GitLab runs **stages sequentially**, but jobs within a stage **concurrently**.
- In the example:

1. All jobs in build must finish before test starts.
 2. All jobs in test must finish before deploy starts.
- This ensures order where needed but allows concurrency to maximize speed.

How Runners Enable Concurrency

- When jobs are triggered, the **GitLab server schedules them** and assigns them to **available runners**.
- If:
 - You have **1 runner** → jobs will run one after another (even in the same stage).
 - You have **multiple runners (or multiple parallel slots in a runner)** → jobs can run **simultaneously**.

What are only and except?

In GitLab CI/CD, every **job** can define **conditions** that decide **when the job should run**.
Two common keywords for this are:

- **only** → Run the job **only if certain conditions are true**.
- **except** → Run the job **except if certain conditions are true**.

These help control **job execution** depending on branch, tag, merge request, or pipeline type.

2. only

You use **only** when you want to **restrict a job to specific cases**.

Examples:

```
deploy_job:
  stage: deploy
  script:
```

```
    - echo "Deploying..."
only:
  - main          # Run only on main branch

tag_build:
  stage: build
  script:
    - echo "Building for release..."
  only:
    - tags          # Run only when a Git tag is pushed

mr_check:
  stage: test
  script:
    - echo "Running checks for MR..."
  only:
    - merge_requests # Run only for merge requests
```

3. except

You use `except` when you want to **exclude some cases**.

Examples:

```
test_job:
  stage: test
  script:
    - echo "Running tests..."
  except:
    - main          # Run everywhere except main branch

dev_build:
  stage: build
  script:
    - echo "Building dev version..."
  except:
    - tags          # Do not run for Git tags
```

4. Combining **only** and **except**

- GitLab will first check **only**.
- If **only** matches, the job is considered for execution.
- Then **except** is checked. If it matches, the job is **excluded**.

Example:

```
special_job:
  stage: test
  script:
    - echo "Special job"
  only:
    - branches
  except:
    - main
```

👉 Meaning: run this job for **all branches except main**.

5. ⚠️ Deprecation Notice

- **only** and **except** are **older syntax**.
- GitLab recommends using **rules:** instead, because rules are **more powerful and flexible**.

Example with rules:

```
deploy_job:
  stage: deploy
  script:
    - echo "Deploying..."
  rules:
```

```
- if: '$CI_COMMIT_BRANCH == "main"'
```

With rules: you can:

- Use conditions (if: with variables like \$CI_COMMIT_BRANCH)
- Control when to run (when: → always, manual, never, on_success, on_failure)
- Combine multiple conditions.

```
# New
```

```
test_job:
```

```
  stage: test
```

```
  script:
```

```
    - echo "Running tests..."
```

```
  rules:
```

```
    - if: '$CI_COMMIT_BRANCH == "main"'
```

```
      when: never # explicitly skip job
```

```
    - when: always # otherwise run
```

```
# New
```

```
tag_build:
```

```
  stage: build
```

```
  script:
```

```
    - echo "Build for release"
```

```
  rules:
```

```
    - if: '$CI_COMMIT_TAG' # variable exists if commit is a tag
```

```
# New
```

```
mr_check:
```

```
  stage: test
```

```
  script:
```

```
    - echo "MR checks"
```

```
  rules:
```

```
    - if: '$CI_PIPELINE_SOURCE == "merge_request_event"'
```

. What are variables in GitLab CI/CD?

- Variables are **key–value pairs** you can use in your `.gitlab-ci.yml` or in your jobs.

Example:

```
job1:  
  stage: test  
  script:  
    - echo "Project name is $CI_PROJECT_NAME"
```

-
- GitLab already provides many **predefined variables** (e.g., `$CI_COMMIT_BRANCH`, `$CI_PROJECT_DIR`).

But you can also **define your own variables**.

2. Defining Variables in `.gitlab-ci.yml`

You can set variables globally or per job.

Global variables (apply to all jobs)

```
variables:  
  APP_ENV: "development"  
  TIMEOUT: "30"
```

Job-specific variables

```
deploy_job:  
  stage: deploy  
  variables:  
    APP_ENV: "production"  
  script:  
    - echo "Deploying in $APP_ENV"
```

3. Setting Variables in GitLab UI (project settings)

This is where **secrets** usually go (like passwords, tokens, API keys).

1. Go to your project in GitLab.
2. Navigate to:
Settings → **CI/CD** → **Variables** → **Expand**.
3. Click “**Add variable**”.
4. Fill in:
 - **Key** → the variable name (e.g., `AWS_ACCESS_KEY_ID`)
 - **Value** → the secret (e.g., your AWS key)
 - Options:
 - **Masked** → hides the value in job logs (recommended for secrets).
 - **Protected** → only available in protected branches/tags (like main or release tags).

After this, the variable is available inside jobs as an **environment variable**.

Example usage in a job:

```
deploy_job:
  stage: deploy
  script:
    - echo "Using AWS key $AWS_ACCESS_KEY_ID"
```

4. Group-level and Instance-level Variables

- **Group variables** → defined once, used by all projects in a group.
- **Instance variables** → set by GitLab admin, available for all projects in an instance.

This avoids duplicating secrets across projects.

5. Secret Files

Sometimes secrets are too large for simple variables (like JSON service accounts, SSH keys, etc.).

You can store them as **file variables**:

- In GitLab UI, check “**Type: File**” when adding a variable.
- In your job, GitLab will **create a temporary file** with that content, and set the variable to the file path.

Example:

```
deploy_job:
  stage: deploy
  script:
    - cat "$GOOGLE_AUTH_JSON"    # prints contents of secret file
```

6. Best Practices

- **Never hardcode secrets** in `.gitlab-ci.yml`.
 - Use **masked + protected variables** for sensitive data.
 - Use **file variables** for large configs.
 - Restrict variable scope if only needed for certain environments (e.g., production).
-

✅ In summary:

- Use `variables:` in `.gitlab-ci.yml` for non-sensitive configs.
- Use GitLab **Settings** → **CI/CD** → **Variables** for secrets.
- Mark them as **masked + protected**.

- Access them in jobs with \$VAR_NAME.

What is when: in GitLab CI/CD?

In your `.gitlab-ci.yml`, each job can have a **when:** instruction.
It tells GitLab **when the job should run** (depending on pipeline conditions).

♦ Common when: options

1. **when: on_success (default)**

👉 Run the job only if all previous jobs succeed.

test_job:

```
script: echo "Running tests"
```

```
when: on_success
```

💡 If you don't write `when: . . .`, this is the default.

2. **when: on_failure**

👉 Run the job only if a previous job fails.

notify_failure:

```
script: echo "Tests failed, sending alert"
```

```
when: on_failure
```

3. **when: always**

👉 Run the job whether previous jobs succeed or fail.

cleanup:

```
script: echo "Cleaning up temp files"
```

```
when: always
```

4. **when: manual**

👉 Job will **not run automatically**. A person must click "Play" in the GitLab UI.

deploy_to_staging:

```
script: echo "Deploying to staging"
```

```
when: manual
```

💡 Useful when you want control over deployment.

5. **when: delayed**

👉 Run the job after a set time delay.

delayed_job:

```
script: echo "Running after 30 seconds"
```

```
when: delayed
```

```
start_in: 30 seconds
```

♦ **Example putting it all together**

stages:

- build
- test
- deploy
- notify

build_job:

stage: build

script: echo "Building project"

test_job:

stage: test

script: echo "Running tests"

deploy_job:

stage: deploy

script: echo "Deploying..."

when: manual # run only when you click "Play"

notify_job:

stage: notify

script: echo "Notifying team"

when: on_failure # only runs if something failed

✅ Summary for beginners

- `on_success` → run if previous jobs succeed (default).
- `on_failure` → run if something failed.
- `always` → run no matter what.
- `manual` → run only if you click it.
- `delayed` → run later after some delay.

◆ `allow_failure`

Normally, if a job **fails**, the pipeline also **fails** ❌.

But if you set `allow_failure: true`, the job may fail, and the pipeline will still continue ✅.

👉 Example:

```
lint:
```

```
  script: echo "Running linting... but may fail"
```

```
  allow_failure: true
```

- If `lint` fails, the pipeline **does not stop**.
- Useful for **optional jobs** (e.g., experimental checks, non-blocking tests).

◆ Other important control flags/options

1. `retry`

Sometimes jobs fail randomly (e.g., network issues). You can tell GitLab to retry automatically.

```
integration_test:
```

```
  script: run-integration-tests.sh
```

```
  retry: 2    # retry up to 2 times if it fails
```

2. timeout

You can set how long a job should run before being killed.

```
long_task:
```

```
  script: do-heavy-processing.sh
```

```
  timeout: 30m    # 30 minutes
```

3. interruptible

Cancels running jobs if a new pipeline is triggered (saves resources).

```
build:
```

```
  script: make build
```

```
  interruptible: true
```

💡 Example: if you push a new commit, the old build is canceled automatically.

4. parallel

Runs multiple instances of the same job in parallel (for faster pipelines).

```
test:
```

```
script: run-tests.sh
```

```
parallel: 5    # split into 5 parallel jobs
```

5. resource_group

Prevents multiple jobs from the same group running at the same time (avoids conflicts).

deploy:

```
script: deploy.sh
```

```
resource_group: production
```

💡 Useful to ensure **only one deploy** to production happens at a time.

◆ Summary (easy view)

- **allow_failure: true** → job may fail, pipeline continues.
- **retry: N** → auto retry job if it fails.
- **timeout: 30m** → stop job if it runs too long.
- **interruptible: true** → cancel old jobs when new ones start.
- **parallel: N** → run multiple copies of a job.
- **resource_group:** → make jobs run one at a time in a group.

Example: tests in parallel

test:

```
script: run-tests.sh
```

```
parallel: 3
```

When the pipeline runs:

- GitLab creates **3 jobs**: test 1/3, test 2/3, test 3/3.
- They all run the same script (run-tests.sh) **in parallel**.
- Useful for splitting heavy work.

◆ Basic: parallel + rules

You just define `rules`: normally, and inside the job you can still say `parallel: N`.

Example: run tests only on main branch, and split into 3 jobs.

```
test:
```

```
  script: run-tests.sh
```

```
  parallel: 3
```

```
  rules:
```

```
    - if: '$CI_COMMIT_BRANCH == "main"'
```

```
      when: always
```

```
    - when: never    # don't run on other branches
```

👉 On main → creates 3 parallel jobs.

👉 On other branches → no job at all.

◆ Smarter: parallel:matrix with rules

You can also make **matrix jobs** with rules.

Example: run tests on multiple Python versions, but only if changes touch backend/ folder.

test:

script: pytest

parallel:

matrix:

- PYTHON: ["3.8", "3.9", "3.10"]

rules:

- changes:

- backend/**/* # only if files in backend changed

👉 If backend code changes → GitLab spawns 3 jobs (PYTHON=3.8, 3.9, 3.10).

👉 If only frontend changes → no jobs created.

◆ Advanced: conditional matrix

You can even control **matrix values with rules**.

Example:

- On main branch → run tests in 3 Python versions.
- On feature branches → only latest version.

test:

script: pytest

parallel:

matrix:


```
- PYTHON: ["3.10"]
```

rules:

```
- if: '$CI_COMMIT_BRANCH == "main"'
```

variables:

```
PYTHON: "3.8,3.9,3.10" # override matrix values
```

1. What is Cache in GitLab CI/CD?

- In GitLab pipelines, each job usually runs in a **fresh environment** (new container, VM, or runner).
- That means any files created in one job are **lost** after the job finishes.
- **Cache** lets you **store files** so they can be **reused by later jobs or future pipeline runs**.

👉 Think of it like a **temporary storage** that speeds things up by avoiding re-downloading or re-installing things every time.

2. What is Cache Used For?

Common use cases:

- Dependencies (e.g., node_modules/, .m2/ for Maven, vendor/ in PHP).
- Build tools downloads (e.g., Gradle wrapper, pip cache).

- Large libraries or assets that don't change often.

It helps pipelines run **faster** and **use fewer resources**.

3. Basic Example

Here's a Node.js example:

stages:

- install
- test

install_dependencies:

stage: install

script:

- npm install

cache:

paths:

- node_modules/

run_tests:

stage: test

script:

- npm test

cache:

paths:

- node_modules/

Explanation:

- First job (install_dependencies) runs npm install and caches node_modules/.
- Second job (run_tests) **reuses the cached node_modules/**, so it doesn't reinstall packages.

4. Cache vs Artifacts (Important!)

- **Cache** → for speeding up jobs across stages or pipelines (dependencies, downloads).
- **Artifacts** → for saving **job outputs** (build results, test reports, binaries) that you want to **access later or download**.

👉 Cache is about **reusing**; artifacts are about **keeping results**.

5. Cache Key

You can define a **cache key** to control when a cache should be reused.

Static key (always same cache)

cache:

key: common-cache

paths:

- node_modules/

Per-branch cache

cache:

key: "\$CI_COMMIT_REF_NAME"

paths:

- node_modules/

👉 Different cache for each branch.

Automatic key (based on files)

cache:

key:

files:

- package-lock.json

paths:

- node_modules/

👉 If package-lock.json changes, a new cache is created.

6. Cache Policy

You can control cache download/upload behavior with policy:

- pull-push (default) → download cache at start, update at end.
- pull → only download cache.
- push → only upload cache.

Example:

cache:

key: my-cache

paths:

- node_modules/

policy: pull # don't overwrite cache, only use existing

7. Cache Limitations

- Cache is stored on the **runner's cache storage** (shared runners may not always share cache between projects).
- Max cache size is **5 GB per project** by default.
- Cache is not guaranteed to persist forever (GitLab may clean old caches).

✅ In summary:

- Cache in GitLab CI/CD = temporary storage for speeding up jobs.
- Useful for dependencies and downloads.
- Defined with `cache: paths: [...]`.
- Controlled with key and policy.
- Different from **artifacts**, which are for saving outputs.

1. The Challenge

By default, GitLab CI/CD jobs run **inside containers** (if using Docker executor).
But:

- You **can't run Docker-in-Docker (dind)** commands directly unless configured.
 - So we need a way to let the pipeline **use Docker** to build images.
-

2. Two Main Approaches

A. Docker-in-Docker (DinD) (most common)

You use a **Docker service container** (`docker:dind`) along with the job.

- `docker:dind` runs the Docker daemon.
- Your job uses the Docker CLI to talk to that daemon.

B. Docker socket binding (faster, but less secure)

You mount the host's `/var/run/docker.sock` into the job container.

- The job uses the **host's Docker daemon** directly.

- Faster, but jobs get host-level privileges.

3. Example: Docker-in-Docker (Recommended by GitLab)

Here's a `.gitlab-ci.yml` that builds and pushes a Docker image:

```
stages:
  - build

build_image:
  stage: build
  image: docker:24.0          # Docker CLI client
  services:
    - docker:24.0-dind        # Docker daemon (DinD)
  variables:
    DOCKER_HOST: tcp://docker:2375/
    DOCKER_TLS_CERTDIR: ""     # Disable TLS (simplifies setup)
  script:
    - docker version
    - docker build -t $CI_REGISTRY_IMAGE:$CI_COMMIT_SHORT_SHA .
    - docker login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD
$CI_REGISTRY
    - docker push $CI_REGISTRY_IMAGE:$CI_COMMIT_SHORT_SHA
```

Explanation:

- `image: docker:24.0` → uses Docker CLI.
- `services: docker:24.0-dind` → starts a Docker daemon container.
- `DOCKER_HOST: tcp://docker:2375/` → tells CLI where to find daemon.
- `$CI_REGISTRY_*` → predefined GitLab variables for container registry.
- `docker build ...` → builds image.

- `docker push ...` → pushes to GitLab Container Registry.

4. Example: Docker Socket Binding (Faster but Risky)

This works only with **GitLab Runner on your own server** (not shared runners).

```
build_image:
  stage: build
  image: docker:24.0
  variables:
    DOCKER_HOST: unix:///var/run/docker.sock
  script:
    - docker build -t myapp:$CI_COMMIT_SHORT_SHA .
    - docker push myapp:$CI_COMMIT_SHORT_SHA
```

Here you mount `/var/run/docker.sock` into the job in runner config. The job talks directly to the host's Docker daemon.

5. Pushing to GitLab Container Registry

Each GitLab project has its own registry.

Registry URL:

`$CI_REGISTRY/$CI_PROJECT_PATH`

For example:

`registry.gitlab.com/mygroup/myproject`

Login & push:

```
docker login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD
$CI_REGISTRY
docker build -t $CI_REGISTRY_IMAGE:$CI_COMMIT_SHORT_SHA .
```



```
docker push $CI_REGISTRY_IMAGE:$CI_COMMIT_SHORT_SHA
```

✓ Summary:

- Use **docker:dind** service in CI jobs to build Docker images.
- Or use **socket binding** if you control the runner.
- Push images to **GitLab Container Registry** using built-in `$CI_REGISTRY_*` variables.

A. Shared Runner

- Provided by GitLab itself (on GitLab.com SaaS) or by an admin in self-hosted GitLab.
- Available to **all projects** in the GitLab instance.
- Good for small projects or when you don't want to manage infrastructure.
- Downside → limited concurrency, slower, no customization (everyone shares the same pool).

👉 Example: If you create a new project on GitLab.com and run a pipeline without installing anything, it uses a **shared runner**.

B. Group Runner

- Installed and registered by you (or your admin) at the **group level**.
- Shared among **all projects in a GitLab group**.
- Useful if you have multiple related projects and want consistent CI/CD setup (same tools, same environment).

👉 Example: You create a group called `my-company/` with multiple repos (frontend, backend, infra) and register one runner for all.

C. Specific Runner

- Registered for **one project only**.
- Fully under your control (tools, resources, permissions).
- Best for projects that need special setup (e.g., CUDA/GPU builds, custom dependencies, access to internal infra).

👉 Example: You have a project that needs **NVIDIA GPU drivers** for ML pipelines, so you install a specific runner on a GPU server only for that project.

D. Kubernetes Executor

- One of the **executor types** GitLab Runner can use.
- Instead of running jobs on a VM or Docker container directly, it spins up a **temporary Pod in Kubernetes** for each CI/CD job.
- Benefits:
 - Scales automatically (each job = new pod).
 - Uses Kubernetes resources efficiently.
 - Isolated and secure (jobs run in separate pods).
- Downsides: needs a working Kubernetes cluster + more setup.

👉 Example: In a large team, you integrate GitLab Runner with your company's Kubernetes cluster → each pipeline job runs in a fresh pod.

A. Standard Variables

- Basic **key-value pairs** you define.

- Can be set:
 - In `.gitlab-ci.yml` (non-secret config)
 - Or in **Project** → **Settings** → **CI/CD** → **Variables**
- Available in **all jobs** by default.

👉 Example (in `.gitlab-ci.yml`):

```
variables:  
  APP_ENV: "dev"
```

Usage:

```
script:  
  - echo "Running in $APP_ENV"
```

B. Group Variables

- Defined at the **group level** (instead of project).
- Automatically inherited by **all projects** in that group.
- Useful when multiple projects need the same config (like cloud keys or repo URLs).

👉 Example:

- You have a GitLab group `mycompany/` with projects (`frontend`, `backend`).
- You set `AWS_REGION=ap-south-1` as a **group variable**.
- Both projects can now use `$AWS_REGION` without redefining it.

C. Protected Variables

- Variables marked as **protected** can only be accessed in pipelines that run on **protected branches/tags** (like main, production).
- Prevents secrets (like prod DB password) from leaking in dev branches or forks.

👉 Example:

- DB_PASSWORD=supersecret is **protected**.
- A pipeline on main can access it.
- A pipeline on feature-xyz **cannot**.

This keeps production secrets safe.

D. Masked Variables

- Secrets that should **never appear in job logs**.
- If a masked variable is echoed accidentally, it shows as ****.
- Must follow certain rules:
 - At least 8 characters long.
 - Only a-z A-Z 0-9 _ = - / + allowed.

👉 Example:

- API_KEY=ABC123XYZ is masked.

If a script does echo \$API_KEY, logs show:

-

✓ Quick Comparison

Type	Where Used	Purpose
Standard Variables	Project or <code>.gitlab-ci.yml</code>	General configs (env, version, timeout)
Group Variables	Group level	Share configs across projects
Protected Variables	Project/Group	Only available in protected branches/tags
Masked Variables	Project/Group	Hide secrets from logs

1. Problem Before needs :

- By default, GitLab pipelines run:
 - **Stages sequentially** (build → test → deploy).
 - **Jobs inside a stage concurrently**.

This means a job in stage **B** can only start after **all jobs** in stage **A** finish, even if it only depends on one of them.

👉 That can waste time.

2. What is needs : ?

- The **needs :** keyword lets you define **direct job-to-job dependencies**, instead of being forced to wait for the whole previous stage.
- This allows **parallel execution across stages** and speeds up pipelines.

👉 Think of it as:

*"This job **needs only X**, not everything before it."*

3. Example Without needs

```
stages:
  - build
  - test
  - deploy

build_job:
  stage: build
  script: echo "Building..."

test_job1:
  stage: test
  script: echo "Unit tests"

test_job2:
  stage: test
  script: echo "Integration tests"

deploy_job:
  stage: deploy
  script: echo "Deploying..."
```

Execution order:

1. build_job runs.
2. After it finishes, **both test_job1 and test_job2** run in parallel.
3. Only after both tests finish → deploy_job starts.

Even if deploy_job only needs test_job1, it still waits for test_job2.

4. Example With needs

```
stages:
  - build
```

- test
- deploy

```
build_job:  
  stage: build  
  script: echo "Building..."
```

```
test_job1:  
  stage: test  
  script: echo "Unit tests"
```

```
test_job2:  
  stage: test  
  script: echo "Integration tests"
```

```
deploy_job:  
  stage: deploy  
  needs: ["test_job1"]  # depend only on test_job1  
  script: echo "Deploying..."
```

Execution order:

- build_job runs.
- Then test_job1 and test_job2 run in parallel.
- As soon as test_job1 finishes, **deploy_job starts** (no need to wait for test_job2).

👉 This shortens pipeline time.

5. Passing Artifacts with needs

- Normally, jobs in later stages can access artifacts from earlier stages.
- With needs :, you can also **download artifacts from specific jobs immediately**.

Example:

```
build_job:
  stage: build
  script: echo "Build artifact" > build.txt
  artifacts:
    paths: [build.txt]

deploy_job:
  stage: deploy
  needs: ["build_job"]    # also fetches build.txt
  script:
    - cat build.txt
    - echo "Deploying..."
```

6. Special Notes

- If you use `needs: []` (empty array), the job starts **immediately** without waiting for its stage.
 - You can mix normal stage order and `needs: overrides`.
 - Be careful: if you misuse it, you can create jobs that start too early and fail (because dependencies aren't ready).
-

Summary:

- `needs:` lets you declare **job dependencies** instead of relying only on stage order.
- Jobs can start as soon as their dependencies finish, **even if other jobs in earlier stages are still running**.
- This makes pipelines **faster** and more **efficient**.